

1. Neural Network Implementation

program

```
import numpy as np
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.metrics import accuracy_score

iris = load_iris()
X = iris.data
y = iris.target.reshape(-1, 1)
encoder = OneHotEncoder(sparse=False)
y = encoder.fit_transform(y)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
Department of Artificial Intelligence and Data Science
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
input_size = X_train.shape[1]
hidden_size = 10
output_size = y_train.shape[1]
learning_rate = 0.01
epochs = 1000
W1 = np.random.randn(input_size, hidden_size)
b1 = np.zeros((1, hidden_size))
W2 = np.random.randn(hidden_size, output_size)
b2 = np.zeros((1, output_size))
def sigmoid(x):
    return 1 / (1 + np.exp(-x))
def sigmoid_derivative(x):
    return x * (1 - x)
for epoch in range(epochs):
    z1 = np.dot(X_train, W1) + b1
    a1 = sigmoid(z1)
    z2 = np.dot(a1, W2) + b2
    a2 = sigmoid(z2)
    loss = np.mean((a2 - y_train) ** 2)
    d_loss_a2 = 2 * (a2 - y_train) / y_train.shape[0]
    d_a2_z2 = sigmoid_derivative(a2)
    d_z2_W2 = a1
    d_z2_a1 = W2
    d_z2 = d_loss_a2 * d_a2_z2
    d_W2 = np.dot(d_z2_W2.T, d_z2)
    d_b2 = np.sum(d_z2, axis=0, keepdims=True)
    d_a1_z1 = sigmoid_derivative(a1)
    d_z1_W1 = X_train
    d_z1 = np.dot(d_z2, d_z2_a1.T) * d_a1_z1
    d_W1 = np.dot(d_z1_W1.T, d_z1)
    d_b1 = np.sum(d_z1, axis=0, keepdims=True)
    W2 -= learning_rate * d_W2
```

```

b2 -= learning_rate * d_b2
W1 -= learning_rate * d_W1
b1 -= learning_rate * d_b1
if epoch % 100 == 0:
    print(f'Epoch {epoch}, Loss: {loss}')
z1 = np.dot(X_test, W1) + b1
a1 = sigmoid(z1)
z2 = np.dot(a1, W2) + b2
a2 = sigmoid(z2)
predictions = np.argmax(a2, axis=1)
y_test_labels = np.argmax(y_test, axis=1)
accuracy = accuracy_score(y_test_labels, predictions)
print(f'Accuracy: {accuracy * 100:.2f}%',)

```

2. TensorFlow Basics

program

```

import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.utils import to_categorical
from sklearn.metrics import accuracy_score
(X_train, y_train), (X_test, y_test) = mnist.load_data()
X_train = X_train / 255.0
X_test = X_test / 255.0
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)
model = Sequential([
    Flatten(input_shape=(28, 28)),
    Dense(128, activation='relu'),
    Dense(64, activation='relu'),
    Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy'])
model.fit(X_train, y_train, epochs=10, batch_size=32,
    validation_split=0.2)
test_loss, test_accuracy = model.evaluate(X_test, y_test)
print(f'Test accuracy: {test_accuracy * 100:.2f}%',)
y_pred = model.predict(X_test)
y_pred_classes = tf.argmax(y_pred, axis=1)
y_true = tf.argmax(y_test, axis=1)
accuracy = accuracy_score(y_true, y_pred_classes)
print(f'Accuracy calculated using sklearn: {accuracy * 100:.2f}%',)
Output

```

3. Building a CNN

program

```
import tensorflow as tf
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import EarlyStopping
(X_train, y_train), (X_test, y_test) = cifar10.load_data()
X_train = X_train / 255.0
X_test = X_test / 255.0
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)
def build_and_compile_model(filter_size, stride, padding):
    model = Sequential([
        Conv2D(32, kernel_size=filter_size, strides=stride, padding=padding,
activation='relu', input_shape=(32, 32, 3)),
        MaxPooling2D(pool_size=(2, 2)),
        Conv2D(64, kernel_size=filter_size, strides=stride, padding=padding,
activation='relu'),
        MaxPooling2D(pool_size=(2, 2)),
        Conv2D(128, kernel_size=filter_size, strides=stride, padding=padding,
activation='relu'),
        Flatten(),
        Dense(128, activation='relu'),
        Dense(10, activation='softmax')
    ])

    model.compile(optimizer='adam',
loss='categorical_crossentropy',
metrics=['accuracy'])
    return model
filter_sizes = [(3, 3), (5, 5)]
strides = [(1, 1), (2, 2)]
paddings = ['valid', 'same']
for filter_size in filter_sizes:
    for stride in strides:
        for padding in paddings:
            print(f"\nExperimenting with filter_size={filter_size}, stride={stride},
padding={padding}")

    model = build_and_compile_model(filter_size, stride, padding)

    early_stopping = EarlyStopping(monitor='val_loss', patience=3)
    history = model.fit(X_train, y_train, epochs=20, batch_size=64,
validation_split=0.2, callbacks=[early_stopping], verbose=1)

    test_loss, test_accuracy = model.evaluate(X_test, y_test)
    print(f"Test accuracy: {test_accuracy * 100:.2f}%")
    model.summary()
```

4. Improving CNN Performance

program

```
import tensorflow as tf
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
Dropout, BatchNormalization
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import EarlyStopping
(X_train, y_train), (X_test, y_test) = cifar10.load_data()
X_train = X_train / 255.0
X_test = X_test / 255.0
y_train = to_categorical(y_train, 10)
Department of Artificial Intelligence and Data Science
y_test = to_categorical(y_test, 10)
datagen = ImageDataGenerator(
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)
datagen.fit(X_train)
def build_and_compile_model():
    model = Sequential([
        Conv2D(32, (3, 3), padding='same', activation='relu', input_shape=(32, 32, 3)),
        BatchNormalization(),
        MaxPooling2D((2, 2)),
        Dropout(0.25),
        Conv2D(64, (3, 3), padding='same', activation='relu'),
        BatchNormalization(),
        MaxPooling2D((2, 2)),
        Dropout(0.25),
        Conv2D(128, (3, 3), padding='same', activation='relu'),
        BatchNormalization(),
        MaxPooling2D((2, 2)),
        Dropout(0.25),
        Flatten(),
        Dense(512, activation='relu'),
        Dropout(0.5),
        Dense(10, activation='softmax')
    ])
    model.compile(optimizer='adam',
        loss='categorical_crossentropy',
        metrics=['accuracy'])
    return model
model = build_and_compile_model()
```

```

early_stopping = EarlyStopping(monitor='val_loss', patience=5,
restore_best_weights=True)
history = model.fit(datagen.flow(X_train, y_train, batch_size=64),
epochs=50,
validation_data=(X_test, y_test),
callbacks=[early_stopping],
verbose=1)
test_loss, test_accuracy = model.evaluate(X_test, y_test)
print(f'Test accuracy: {test_accuracy * 100:.2f}%')
import matplotlib.pyplot as plt
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(['Train', 'Validation'])
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
Department of Artificial Intelligence and Data Science
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend(['Train', 'Validation'])
plt.show()

```

5. RNN for sequence prediction using TensorFlow

program

```

import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import SimpleRNN, Dense
from sklearn.preprocessing import MinMaxScaler
import matplotlib.pyplot as plt
def generate_sine_wave(seq_length, num_samples):
    x = np.linspace(0, 2 * np.pi, seq_length)
    Department of Artificial Intelligence and Data Science
    data = np.sin(x)
    sequences = []
    for i in range(num_samples):
        start = np.random.randint(0, len(data) - seq_length)
        sequences.append(data[start:start + seq_length])
    return np.array(sequences)
seq_length = 50
num_samples = 1000
num_epochs = 10
batch_size = 32
data = generate_sine_wave(seq_length, num_samples)

```

```

X = data[:, :-1]
y = data[:, -1]
scaler = MinMaxScaler()
X = scaler.fit_transform(X)
y = scaler.transform(y.reshape(-1, 1)).flatten()
X = X.reshape((num_samples, seq_length - 1, 1))
model = Sequential([
    SimpleRNN(50, activation='relu', input_shape=(X.shape[1], 1)),
    Dense(1)
])
model.compile(optimizer='adam', loss='mse')
model.fit(X, y, epochs=num_epochs, batch_size=batch_size)
predictions = model.predict(X)
y = scaler.inverse_transform(y.reshape(-1, 1)).flatten()
predictions = scaler.inverse_transform(predictions).flatten()
plt.figure(figsize=(12, 6))
plt.plot(range(len(y)), y, label='True Values')
Department of Artificial Intelligence and Data Science
plt.plot(range(len(predictions)), predictions, label='Predicted Values')
plt.legend()
plt.show()
Output

```

6. LSTM network for text generation

program

```

import numpy as np
import tensorflow as tf
Department of Artificial Intelligence and Data Science
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.sequence import pad_sequences
# Load and preprocess the text data
def load_text(file_path):
    with open(file_path, 'r') as file:
        text = file.read().lower()
    return text
def prepare_data(text, seq_length):
    tokenizer = Tokenizer(char_level=True)
    tokenizer.fit_on_texts([text])
    total_chars = len(tokenizer.word_index) + 1
    sequences = []
    next_chars = []
    for i in range(0, len(text) - seq_length, 1):
        seq = text[i:i + seq_length]
        label = text[i + seq_length]
        sequences.append([tokenizer.char_index[c] for c in seq])
        next_chars.append(tokenizer.char_index[label])
    X = np.array(sequences)

```

```

y = np.array(next_chars)
X = pad_sequences(X, maxlen=seq_length)
y = to_categorical(y, num_classes=total_chars)
return X, y, tokenizer, total_chars
Department of Artificial Intelligence and Data Science
def build_model(seq_length, total_chars):
    model = Sequential([
        LSTM(128, input_shape=(seq_length, total_chars),
return_sequences=True),
        Dropout(0.2),
        LSTM(128),
        Dense(total_chars, activation='softmax')
    ])
    model.compile(loss='categorical_crossentropy', optimizer='adam',
metrics=['accuracy'])
    return model
def train_model(model, X, y, epochs=20, batch_size=128):
    model.fit(X, y, epochs=epochs, batch_size=batch_size, verbose=1)
def generate_text(model, tokenizer, total_chars, seed_text, seq_length,
num_chars):
    reverse_char_index = {i: char for char, i in tokenizer.char_index.items()}
    generated_text = seed_text
    for _ in range(num_chars):
        x_pred = pad_sequences([tokenizer.texts_to_sequences([seed_text])[0]],
maxlen=seq_length)
        pred_probs = model.predict(x_pred, verbose=0)[0]
        next_index = np.argmax(pred_probs)
        next_char = reverse_char_index[next_index]
        generated_text += next_char
        seed_text = seed_text[1:] + next_char
    return generated_text
def main():
    text = load_text('shakespeare.txt')
    seq_length = 40
Department of Artificial Intelligence and Data Science
    X, y, tokenizer, total_chars = prepare_data(text, seq_length)
    model = build_model(seq_length, total_chars)
    train_model(model, X, y, epochs=20, batch_size=128)
    seed_text = text[:seq_length]
    generated_text = generate_text(model, tokenizer, total_chars, seed_text,
seq_length, num_chars=400)
    print(generated_text)
if __name__ == "__main__":
    main()

```

Output

Department of Artificial Intelligence and Data Science

7. Q-learning algorithm to solve the FrozenLake environment

program

```
import gym
import numpy as np
env = gym.make('FrozenLake-v1', is_slippery=True, render_mode=None)
learning_rate = 0.1
discount_factor = 0.99
epsilon = 1.0 # Exploration rate
max_epsilon = 1.0
min_epsilon = 0.01
decay_rate = 0.005 # Decay rate for exploration probability
state_size = env.observation_space.n
action_size = env.action_space.n
Department of Artificial Intelligence and Data Science
q_table = np.zeros((state_size, action_size))
num_episodes = 10000
max_steps = 100
for episode in range(num_episodes):
    state = env.reset()[0]
    done = False
    for step in range(max_steps):
        if np.random.uniform(0, 1) < epsilon:
            action = env.action_space.sample() # Explore
        else:
            action = np.argmax(q_table[state, :]) # Exploit
        new_state, reward, done, truncated, info = env.step(action)
        q_table[state, action] = q_table[state, action] + learning_rate * (
            reward + discount_factor * np.max(q_table[new_state, :]) -
            q_table[state, action])
        state = new_state
        if done:
            break
    epsilon = min_epsilon + (max_epsilon - min_epsilon) * np.exp(-decay_rate *
episode)
num_test_episodes = 100
total_rewards = 0
for episode in range(num_test_episodes):
    state = env.reset()[0]
    done = False
    episode_rewards = 0
    for step in range(max_steps):
        action = np.argmax(q_table[state, :]) # Choose best action
        Department of Artificial Intelligence and Data Science
        new_state, reward, done, truncated, info = env.step(action)
        episode_rewards += reward
        state = new_state
        if done:
            break
    total_rewards += episode_rewards
print(f"Average reward over {num_test_episodes} test episodes: {total_rewards
/ num_test_episodes}")
```


env.close()
Output
Department of Artificial Intelligence and Data Science

8. Q-values in a Deep Q-Network (DQN)

program

```
import gym
import numpy as np
import random
from collections import deque
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam
EPISODES = 1000
STATE_SIZE = 4
ACTION_SIZE = 2
GAMMA = 0.95
EPSILON = 1.0
EPSILON_MIN = 0.01
EPSILON_DECAY = 0.995
LEARNING_RATE = 0.001
BATCH_SIZE = 32
MEMORY_SIZE = 2000
TARGET_UPDATE = 10
class DQNAgent:
    def __init__(self):
        self.memory = deque(maxlen=MEMORY_SIZE)
        self.epsilon = EPSILON
        self.model = self._build_model()
        self.target_model = self._build_model()
        self.update_target_model()
    def _build_model(self):
        model = Sequential()
        model.add(Dense(24, input_dim=STATE_SIZE, activation='relu'))
        model.add(Dense(24, activation='relu'))
        model.add(Dense(ACTION_SIZE, activation='linear'))
        model.compile(loss='mse',
            optimizer=Adam(learning_rate=LEARNING_RATE))
        return model
    def update_target_model(self):
        self.target_model.set_weights(self.model.get_weights())
    def remember(self, state, action, reward, next_state, done):
        self.memory.append((state, action, reward, next_state, done))
    def act(self, state):
        if np.random.rand() <= self.epsilon:
            return random.randrange(ACTION_SIZE)
        q_values = self.model.predict(state)
        return np.argmax(q_values[0])
    def replay(self):
```

```

if len(self.memory) < BATCH_SIZE:
    return
minibatch = random.sample(self.memory, BATCH_SIZE)
for state, action, reward, next_state, done in minibatch:
    target = self.model.predict(state)
    if done:
        target[0][action] = reward
    else:
        t = self.target_model.predict(next_state)
        target[0][action] = reward + GAMMA * np.amax(t[0])
    self.model.fit(state, target, epochs=1, verbose=0)
    if self.epsilon > EPSILON_MIN:
        self.epsilon *= EPSILON_DECAY
def load(self, name):
    self.model.load_weights(name)
def save(self, name):
    self.model.save_weights(name)
if __name__ == "__main__":
    env = gym.make('CartPole-v1')
    agent = DQNAgent()
    done = False
    for e in range(EPISODES):
        state = env.reset()
        state = np.reshape(state, [1, STATE_SIZE])
        for time in range(500):
            action = agent.act(state)
            next_state, reward, done, _ = env.step(action)
            reward = reward if not done else -10
            next_state = np.reshape(next_state, [1, STATE_SIZE])
            agent.remember(state, action, reward, next_state, done)
            state = next_state
        if done:
            agent.update_target_model()
            print(f'Episode: {e}/{EPISODES}, score: {time}, epsilon: {agent.epsilon:.2}')
            break
    agent.replay()

```

9. Policy gradient method (REINFORCE) to solve the LunarLander environment

program

```

import gym
import numpy as np
Department of Artificial Intelligence and Data Science
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam
EPISODES = 1000

```

```

LEARNING_RATE = 0.001
GAMMA = 0.99
class REINFORCEAgent:
    def __init__(self, state_size, action_size):
        self.state_size = state_size
        self.action_size = action_size
        self.model = self._build_model()
        self.optimizer = Adam(learning_rate=LEARNING_RATE)
    def _build_model(self):
        model = Sequential()
        model.add(Dense(24, input_dim=self.state_size, activation='relu'))
        model.add(Dense(24, activation='relu'))
        model.add(Dense(self.action_size, activation='softmax'))
        return model
    def choose_action(self, state):
        state = state[np.newaxis, :]
        probabilities = self.model.predict(state, verbose=0)[0]
        return np.random.choice(self.action_size, p=probabilities)
    def compute_discounted_rewards(self, rewards):
        discounted_rewards = np.zeros_like(rewards, dtype=np.float32)
        cumulative_reward = 0
        Department of Artificial Intelligence and Data Science
        for t in reversed(range(len(rewards))):
            cumulative_reward = rewards[t] + GAMMA * cumulative_reward
            discounted_rewards[t] = cumulative_reward
        return discounted_rewards
    def train(self, states, actions, discounted_rewards):
        with tf.GradientTape() as tape:
            action_probs = self.model(states, training=True)
            action_indices = tf.range(tf.shape(actions)[0]) * self.action_size + actions
            selected_action_probs = tf.gather(tf.reshape(action_probs, [-1]),
            action_indices)
            loss = -tf.reduce_mean(tf.math.log(selected_action_probs) *
            discounted_rewards)
            gradients = tape.gradient(loss, self.model.trainable_variables)
            self.optimizer.apply_gradients(zip(gradients,
            self.model.trainable_variables))
if __name__ == "__main__":
    env = gym.make('LunarLander-v2')
    state_size = env.observation_space.shape[0]
    action_size = env.action_space.n
    agent = REINFORCEAgent(state_size, action_size)
    for episode in range(EPISODES):
        state = env.reset()
        states, actions, rewards = [], [], []
        while True:
            action = agent.choose_action(state)
            next_state, reward, done, _ = env.step(action)
            states.append(state)
            Department of Artificial Intelligence and Data Science
            actions.append(action)
            rewards.append(reward)

```

```

state = next_state
if done:
    states = np.array(states)
    actions = np.array(actions)
    discounted_rewards = agent.compute_discounted_rewards(rewards)
    discounted_rewards -= np.mean(discounted_rewards)
    discounted_rewards /= np.std(discounted_rewards) + 1e-8
    agent.train(states, actions, discounted_rewards)
    print(f"Episode: {episode + 1}, Reward: {sum(rewards)}")
    break
env.close()

```

10. Imitation learning to train an autonomous vehicle agent

program

```

import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import Adam
from sklearn.model_selection import train_test_split
def load_data():
    num_samples = 10000
    state_dim = 10 # Example state dimensions (e.g., sensor readings)
    action_dim = 3 # Example action dimensions (e.g., steer, throttle, brake)
    states = np.random.rand(num_samples, state_dim)
    actions = np.random.rand(num_samples, action_dim)
    return states, actions
Department of Artificial Intelligence and Data Science
def build_model(state_dim, action_dim):
    model = Sequential()
    model.add(Dense(64, input_dim=state_dim, activation='relu'))
    model.add(Dropout(0.1))
    model.add(Dense(64, activation='relu'))
    model.add(Dense(action_dim, activation='linear'))
    model.compile(optimizer=Adam(learning_rate=0.001), loss='mse')
    return model
if __name__ == "__main__":
    states, actions = load_data()
    states_train, states_val, actions_train, actions_val = train_test_split(states,
actions, test_size=0.2, random_state=42)
    model = build_model(states.shape[1], actions.shape[1])
    model.fit(states_train, actions_train, validation_data=(states_val, actions_val),
epochs=50, batch_size=32)
    loss = model.evaluate(states_val, actions_val)
    print(f"Validation Loss: {loss}")
    new_state = np.random.rand(1, states.shape[1]) # Dummy state
    predicted_action = model.predict(new_state)
    print(f"Predicted Action: {predicted_action}")
Department of Artificial Intelligence and Data Science

```

11.Simplified version of ChaufferNet

program

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, Flatten, Dropout
from tensorflow.keras.optimizers import Adam
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
def load_data():
    num_samples = 1000
    img_height, img_width, img_channels = 64, 64, 3
    Department of Artificial Intelligence and Data Science
    features = np.random.rand(num_samples, img_height, img_width,
    img_channels)
    steering_angles = np.random.rand(num_samples) * 2 - 1 # Range [-1, 1]
    return features, steering_angles
def build_model(input_shape):
    model = Sequential()
    model.add(Conv2D(24, (5, 5), strides=(2, 2), activation='relu',
    input_shape=input_shape))
    model.add(Dropout(0.2))
    model.add(Conv2D(36, (5, 5), strides=(2, 2), activation='relu'))
    model.add(Dropout(0.2))
    model.add(Conv2D(48, (5, 5), strides=(2, 2), activation='relu'))
    model.add(Dropout(0.2))
    model.add(Flatten())
    model.add(Dense(100, activation='relu'))
    model.add(Dropout(0.2))
    model.add(Dense(50, activation='relu'))
    model.add(Dense(10, activation='relu'))
    model.add(Dense(1))
    model.compile(optimizer=Adam(learning_rate=0.001), loss='mse')
    return model
if __name__ == "__main__":
    features, steering_angles = load_data()
    features_train, features_val, angles_train, angles_val =
    train_test_split(features, steering_angles, test_size=0.2, random_state=42)
    input_shape = features_train.shape[1:]
    model = build_model(input_shape)
    model.fit(features_train, angles_train, validation_data=(features_val,
    angles_val), epochs=10, batch_size=32)
    Department of Artificial Intelligence and Data Science
    loss = model.evaluate(features_val, angles_val)
    print(f"Validation Loss: {loss}")
    new_feature = np.random.rand(1, *input_shape)
    predicted_angle = model.predict(new_feature)
    print(f"Predicted Steering Angle: {predicted_angle[0][0]}")
```

12.End-to-End Deep Learning for Autonomous Driving

program

```
import os
import numpy as np
import cv2
import csv
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, Flatten, Dense, Lambda,
Dropout
from tensorflow.keras.optimizers import Adam
from sklearn.model_selection import train_test_split
from sklearn.utils import shuffle
def load_data(data_dir):
    images = []
    steering_angles = []
    with open(os.path.join(data_dir, 'driving_log.csv')) as csvfile:
        reader = csv.reader(csvfile)
        next(reader)
        Department of Artificial Intelligence and Data Science
        for line in reader:
            center_image_path = os.path.join(data_dir, 'IMG', line[0].split('/')[-1])
            center_image = cv2.imread(center_image_path)
            center_image = cv2.cvtColor(center_image, cv2.COLOR_BGR2RGB)
            center_image = cv2.resize(center_image, (200, 66))
            images.append(center_image)
            steering_angles.append(float(line[3])) # Steering angle
    return np.array(images), np.array(steering_angles)
def build_model(input_shape):
    model = Sequential()
    model.add(Lambda(lambda x: x / 255.0 - 0.5,
input_shape=input_shape))
    model.add(Conv2D(24, (5, 5), strides=(2, 2), activation='relu'))
    model.add(Conv2D(36, (5, 5), strides=(2, 2), activation='relu'))
    model.add(Conv2D(48, (5, 5), strides=(2, 2), activation='relu'))
    model.add(Conv2D(64, (3, 3), activation='relu'))
    model.add(Conv2D(64, (3, 3), activation='relu'))
    model.add(Flatten())
    model.add(Dense(100, activation='relu'))
    model.add(Dropout(0.5))
    model.add(Dense(50, activation='relu'))
    model.add(Dense(10, activation='relu'))
    model.add(Dense(1)) # Output layer for steering angle
    model.compile(optimizer=Adam(learning_rate=0.0001), loss='mse')
    return model
if __name__ == "__main__":
    Department of Artificial Intelligence and Data Science
    data_dir = 'path_to_simulator_data' # Replace with your data directory
    images, steering_angles = load_data(data_dir)
```

```
X_train, X_val, y_train, y_val = train_test_split(images,
steering_angles, test_size=0.2, random_state=42)
input_shape = X_train.shape[1:]
model = build_model(input_shape)
model.fit(X_train, y_train, validation_data=(X_val, y_val), epochs=10,
batch_size=32, shuffle=True)
loss = model.evaluate(X_val, y_val)
print(f"Validation Loss: {loss}")
model.save('autonomous_driving_model.h5')
new_image = np.expand_dims(X_val[0], axis=0)
predicted_angle = model.predict(new_image)
print(f"Predicted Steering Angle: {predicted_angle[0][0]}")
```